

```
In [955]: #Submitted by : Philani Mporu (1848751)
#Submitted by : Matthew Kruger (1669326)
#Submitted by : Chloë Smith (1877342)
#Submitted by : Gbolan Fared Bangie (1828281)
```

```
In [ ]: import numpy as np
import pandas as pd
import matplotlib.pyplot as plt
import math
```

```
In [956]: # 4(a) i. Sample 150 x-values from a Normal distribution using a mean of 0 and standard deviation of 10.
# ii. From the x-values construct a design matrix using the features {1,x,x2}.
xValues = np.random.normal(0, 10, 150)
ones = np.ones(150)
ninety = np.empty(90)
x = np.empty(150)
xSquare = np.empty(150)
xCube = np.empty(150)

for i in range(150):
    x[i] = xValues[i]
    xSquare[i] = xValues[i]**2
    xCube[i] = xValues[i]**3
#xCube.sort()
data = pd.DataFrame(ones)
data.columns = ['1']
data['x'] = x
data['x^2'] = xSquare
#print(data)
```

```
In [957]: # iii. Use a uniform distribution to sample true values for  $\theta_0$ ,  $\theta_1$  and  $\theta_2$ .
thetaValues = np.random.uniform(0, 1, 3)
print("True values for theta = " + str(thetaValues))

# iv. Use your design matrix and the true parameters you obtained to create the y-values for the regression data. Finally add random noise to the y-values using a Normal distribution with mean  $\theta_0$  and standard deviation of 8.

yValues = np.empty(150) # Use for when we do Gradient Descent.
yValues4 = np.empty(150) # Use when an extra theta parameter is added.
yValuesReg = np.empty(150) # Use for when we do Gradient Descent with 4 parameters with Regularization.
noiseValues = np.random.normal(0, 8, 150)


for i in range(150):
    yValues[i] = (thetaValues[0] + thetaValues[1] * data['x'][i] + thetaValues[2] * data['x^2'][i]) + noiseValues[i]
    data['y'] = yValues
    np.round(data,6)
    #data.sort_values(by=['x'], inplace=True, ascending=True)
    print(data)

True values for theta = [0.75397704 0.35218594 0.47469017]
   1      x      x^2      y
0  1.0 -0.530356  0.281277  18.991994
1  1.0 -5.809310  33.748079  23.484538
2  1.0  1.816351  3.299130  1.126770
3  1.0 -0.382375  0.146139  19.727272
4  1.0  8.766583  76.852983  29.951098
... ..
145 1.0  1.218888  1.485689  11.773456
146 1.0  33.051451 1092.398382  533.564854
147 1.0 17.136417  293.656793  150.021559
148 1.0 12.195628  148.733354  79.007874
149 1.0 -19.410461 376.765991 162.756669

[150 rows x 4 columns]
```

```
In [958]: # vii. Plot the x-values and their corresponding y-values on a 2D-axis. Your data should look similar to the data shown in Figure 2a. Hint: pyplot.scatter

plt.scatter(data['x'], data['y'])
plt.title('Scatter Plot')
plt.xlabel('x')
plt.ylabel('y')
plt.show()
```



```
In [959]: # vi. Split the data into training, validation and test datasets.

#Training
trainDataX = data['x'][0:90]
trainDataY = data['y'][0:90]
trainData = pd.DataFrame({'x':trainDataX,'y':trainDataY})
#print(trainData)

#Validation
validationDataX = data['x'][90:110]
validationDataY = data['y'][90:110]
validationData = pd.DataFrame({'x':validationDataX,'y':validationDataY})
#print(validationData)

#Testing
testDataX = data['x'][110:150]
testDataY = data['y'][110:150]
testData = pd.DataFrame({'x':testDataX,'y':testDataY})
#print(testData)

# (b) i. Use the Moore-Penrose pseudo-inverse to calculate the closed form solution for the model's parameter values.
# X is the data.drop(columns = 'y') matrix.
# Break up the computation into smaller parts.

X = data.drop(columns = 'y')
multi_inverse = np.linalg.inv(np.dot(X.T,X)) # inverse(X.T dot X)
inverse_dot_XT = np.dot(multi_inverse, X.T) # inverse(X.T dot X) dot X.T
learnValues = np.dot(inverse_dot_XT,data['y']) # (inverse(X.T dot X) dot X.T) dot data ['y']
print("Closed form solution values for theta = " + str(learnValues)) #Moore-Penrose pseudo-inverse

# ii. How close are the learned parameter values to the true parameter values we used to generate the data?

diffValues = [abs(thetaValues[0] - learnValues[0]), abs(thetaValues[1] - learnValues[1]), abs(thetaValues[2] - learnValues[2])]
print("Using the closed form solution, learnt parameter values differ from the true values by the following amount: ")
print(diffValues)

Closed form solution values for theta = [0.53668099 0.36135889 0.47126411]
Using the closed form solution, learnt parameter values differ from the true values by the following amount:
[0.21708984240264912, 0.00914995431517885, 0.0034260587432572986]
```

```
In [960]: # iii. Compute the training error and validation error for the learned regression model.
dataLearn = pd.DataFrame(ones)
dataLearn.columns = ['1']
dataLearn['x'] = x
dataLearn['x^2'] = xSquare

yValuesLearn = np.empty(150)
for i in range(150):
    yValuesLearn[i] = (learnValues[0] + dataLearn['x^2'][i] + learnValues[1] * dataLearn['x'][i] + learnValues[2] * dataLearn['x^2'][i]) + noiseValues[i]
    dataLearn['y'] = yValuesLearn

#Learnt Training
learnTrainData = pd.DataFrame(ninety)
learnTrainData.columns = ['1']
learnTrainData['x'] = dataLearn['x'][0:90]
learnTrainData['y'] = dataLearn['y'][0:90]
dataLearn.sort_values(by=['x'], inplace=True, ascending=True)
#print(learnTrainData)

#Learnt Validation
learnValidationDataX = dataLearn['x'][90:110]
learnValidationDataY = dataLearn['y'][90:110]
learnValidationData = pd.DataFrame({'x':learnValidationDataX,'y':learnValidationDataY})
#print(learnValidationData)

#Computing the training error of closed form solution:
for i in range(90):
    total = 0
    E = (data['y'][i] - dataLearn['y'][i])**2
    total = total + E
trainingErrorCF = 0.5 * total
print("Training Error of closed form solution = " + str(trainingErrorCF))


#Computing the validation error of closed form solution:
for i in range(30):
    total = 0
    E = (data['y'][90 + i] - dataLearn['y'][90 + i])**2
    total = total + E
validationErrorCF = 0.5 * total
print("Validation Error of closed form solution = " + str(validationErrorCF))

# iv. Create a scatter plot of the individual data points along with the learned regression function,
# your plot should look like Figure 2b. Hint: pyplot.plot, this plotting function will give weird
# results if the x-values of the data are not sorted. x train[x train[:,1].argsort()] will give you the
# design matrix for your training data sorted by the second column (where the x values should be).

print(dataLearn)
plt.scatter(dataLearn['x'], dataLearn['y'])
plt.plot(dataLearn['x'], dataLearn['y'], color = 'red')
plt.title('Line Plot')
plt.xlabel('x')
plt.ylabel('y')
plt.show()

Training Error of closed form solution = 0.022918594175590796
Validation Error of closed form solution = 0.022343089662256663
   1      x      x^2      y
32  1.0  24.299045  590.447608  266.027814
44  1.0  24.157820  583.600281  271.483423
47  1.0 -21.662895  469.281009  204.678884
80  1.0 -20.839179  434.271370  197.727344
127 1.0 20.359215  414.497619 190.293308
... ..
82  1.0 17.996908  323.888687  175.469827
151 1.0 18.264678  331.418319 164.221033
23  1.0 19.946288  397.854389 182.723556
50  1.0 23.765227  564.786014 277.877993
146 1.0 33.051451 1092.398382 529.907644

[150 rows x 4 columns]
```



```
In [961]: # v. Repeat the above process using Gradient Descent to train your model. of closed form solution

thetaGD = np.random.uniform(0, 1, 3)

def predictedY(matrix,thetaGuess, index):
    temp = thetaGuess[0]*matrix[0][index] + thetaGuess[1]*matrix[1][index] + thetaGuess[2]*matrix[2][index]
    #print(str(temp) + " predicted y ")
    return temp

def gradDescent(matrix, yValues):
    thetaGuess = theta = np.random.uniform(0,1,3)
    learn = 0.0000000000000001
    tol = 0.000001
    maxRuns = 0
    # print(matrix[0,5])
    for i in range(129):
        oldGuess = thetaGuess
        temp = (predictedY(matrix, thetaGuess, i) - yValues[i])
        # print(str(yValues[i]) + " actual y")
        # print(temp)
        temp = temp*learn
        temp = temp*np.array([matrix[0][i],matrix[1][i],matrix[2][i]])
        #print(temp)
        thetaGuess = thetaGuess - temp
        # print(str(thetaGuess) + " + 'old'")
        # print(str(thetaGuess) + " + 'guess'")
        if (np.linalg.norm(thetaGuess - oldGuess) < tol) or (maxRuns == 1000):
            break
    else:
        maxRuns = maxRuns + 1
        if i == 129:
            i = 0
    # print(thetaGuess)
    return thetaGuess

matrix = X.to_numpy() # Convert DataFrame to numpy
y = data['y'].to_numpy() # Convert DataFrame to numpy
to make it work
thetaGuessGD = gradDescent(matrix, y) # Find values for theta using Gradient Descent.
print("True values for theta = " + str(thetaValues))
print("Values using Gradient Descent = " + str(thetaGuessGD))

# In addition, plot the training error of your regression model over time (observe or capture the training error every 20 parameter updates/time steps). Your plot should look like Figure 2c.

trainingErrorArray = np.zeros(90) # Used to track the error of the model over time.
for i in range(150):
    yValuesGD[i] = (thetaGuessGD[0] + thetaGuessGD[1] * data['x'][i] + thetaGuessGD[2] * data['x^2'][i]) + noiseValues[i]
    data['yGD'] = yValuesGD
print(" ")

#Computing the training error of Gradient of Descent and plot it on a graph:
for i in range(90):
    total = 0
    E = (data['y'][i] - data['yGD'][i])**2
    total = total + E
trainingErrorGD = 0.5 * total
print("Training Error for Gradient Descent = " + str(trainingErrorGD))

#Computing the validation error of Gradient Descent:
for i in range(30):
    total = 0
    E = (data['y'][90 + i] - data['yGD'][90 + i])**2
    total = total + E
validationErrorGD = 0.5 * total
print("Validation Error for Gradient Descent = " + str(validationErrorGD))


#Computing the test error of Gradient Descent:
for i in range(30):
    total = 0
    E = (data['y'][110 + i] - data['yGD'][110 + i])**2
    total = total + E
testErrorGD = 0.5 * total
print("Test Error for Gradient Descent = " + str(testErrorGD))

trainingErrorArray[::1].sort()
plt.plot(trainingErrorArray, color = 'orange')
plt.title('Training Error over Time')
plt.xlabel('Timestep')
plt.ylabel('Training Error')
plt.show()

print(data)

True values for theta = [0.75397704 0.35218594 0.47469017]
Values using Gradient Descent = [0.3248695 0.39728912 0.52779145]

Training Error for Gradient Descent = 0.08480904245109342
Validation Error for Gradient Descent = 0.0256841716252859
Test Error for Gradient Descent = 30.415045948047496
```



	1	x	x^2	y	yGD
0	1.0	-0.530356	0.281277	18.991994	18.563449
1	1.0	-5.809310	33.748079	23.484538	24.690946
2	1.0	1.816351	3.299130	1.126770	0.922079
3	1.0	-0.382375	0.146139	19.727272	32.987551
4	1.0	8.766583	76.852983	29.951098	33.848585
...
145	1.0	1.218888	1.485689	11.773456	11.456276
146	1.0	33.051451	1092.398382	533.564854	592.039303
147	1.0	17.136417	293.656793	150.021559	165.650456
148	1.0	12.195628	148.733354	79.007874	86.807239
149	1.0	-19.410461	376.765991	162.756669	181.808234

```
In [962]: # (c) i. Append a third feature to your design matrix for x^3
thetaValues4 = np.random.uniform(0, 1, 4) # Will be used to generate 4 true values for theta.
print("True values for theta = " + str(thetaValues4))
X['x^3'] = xCube
print(X)

True values for theta = [0.75206039 0.33073176 0.48262031 0.54563629]
   1      x      x^2      x^3      y
0  1.0 -0.530356  0.281277  -0.149177
1  1.0 -5.809310  33.748079 -196.053041
2  1.0  1.816351  3.299130  5.992377
3  1.0 -0.382375  0.146139 -0.5691039
4  1.0  8.766583  76.852983  673.738975
... ..
145 1.0  1.218888  1.485689  1.010809
146 1.0 33.051451 1092.398382 36105.351998
147 1.0 17.136417  293.656793 5032.225299
148 1.0 12.195628  148.733354 1813.896728
149 1.0 -19.410461 376.765991 -7313.201532

[150 rows x 5 columns]
```

```
In [963]: # ii. Train a model using Gradient Descent with the new design matrix. Repeat the process used above.
# In Question 4b, Note, we are now using a third-order polynomial to fit data which was generated using a second-order polynomial. Our function is, thus, more complicated than is necessary to fit the data and as a result will overfit

def predictedY4(matrix,thetaGuess, index):
    temp = thetaGuess[0]*matrix[0][index] + thetaGuess[1]*matrix[1][index] + thetaGuess[2]*matrix[2][index] + thetaGuess[3]*matrix[3][index]
    #print(str(temp) + " predicted y ")
    return temp

def gradDescent4(matrix, yValues):
    thetaGuess = theta = np.random.uniform(0,1,4)
    learn = 0.0000000000000001
    tol = 0.000001
    maxRuns = 0
    # print(matrix[0,5])
    for i in range(129):
        oldGuess = thetaGuess
        temp = (predictedY4(matrix, thetaGuess, i) - yValues[i])
        # print(str(yValues[i]) + " actual y")
        #print(temp)
        temp = temp*learn
        temp = temp*np.array([matrix[0][i],matrix[1][i],matrix[2][i],matrix[3][i]])
        #print(temp)
        thetaGuess = thetaGuess - temp
        # print(str(thetaGuess) + " + 'old'")
        # print(str(thetaGuess) + " + 'guess'")
        if (np.linalg.norm(thetaGuess - oldGuess) < tol) or (maxRuns == 1000):
            break
    else:
        maxRuns = maxRuns + 1
        if i == 129:
            i = 0
    # print(thetaGuess)
    return thetaGuess

matrix4 = X.to_numpy() # Convert DataFrame to numpy
y to make it work
thetaGuess4 = gradDescent4(matrix4, y) # Find parameters for theta using Gradient Descent.
print("True values for theta = " + str(thetaValues4))
print("Values for theta using GD = " + str(thetaGuess4))

for i in range(150):
    yValues4[i] = (thetaGuess4[0] + thetaGuess4[1] * X['x'][i] + thetaGuess4[2] * X['x^2'][i] + thetaGuess4[3] * X['x^3'][i]) + noiseValues[i]
print(" ")

#Computing the training error of Gradient of Descent of third order polynomial:
for i in range(90):
    total = 0
    E = (data['y'][i] - yValues4[i])**2
    total = total + E
trainingError4 = 0.5 * total
print("Training Error for Gradient Descent 3rd order = " + str(trainingError4))

#Computing the validation error of Gradient Descent of third order polynomial:
for i in range(30):
    total = 0
    E = (data['y'][90 + i] - yValues4[90 + i])**2
    total = total + E
validationError4 = 0.5 * total
print("Validation Error for Gradient Descent 3rd order = " + str(validationError4))

#Computing the test error of Gradient Descent:
for i in range(30):
    total = 0
    E = (data['y'][110 + i] - yValues4[110 + i])**2
    total = total + E
testErrorGD4 = 0.5 * total
print("Test Error for Gradient Descent 3rd Order = " + str(testErrorGD4))

True values for theta = [0.75206039 0.33073176 0.48262031 0.54563629]
Values for theta using GDReg = [0.72234895 0.65362995 0.79008449 0.99912847]
```

```
In [964]: # iii. Repeat the training process one final time, this time use regularization when training the third order polynomial model.

def predictedYReg(matrix,thetaGuess, index):
    temp = thetaGuess[0]*matrix[0][index] + thetaGuess[1]*matrix[1][index] + thetaGuess[2]*matrix[2][index] + thetaGuess[3]*matrix[3][index]
    #print(str(temp) + " predicted y ")
    return temp

def gradDescentReg(matrix, yValues):
    thetaGuess = theta = np.random.uniform(0,1,4)
    learn = 0.0000000000000001
    tol = 0.000001
    maxRuns = 0
    # print(matrix[0,5])
    for i in range(129):
        oldGuess = thetaGuess
        temp = (predictedYReg(matrix, thetaGuess, i) - yValues[i])
        # print(str(yValues[i]) + " actual y")
        #print(temp)
        temp = temp*learn
        if i == 0:
            temp = temp*np.array([matrix[0][i],matrix[1][i],matrix[2][i],matrix[3][i]])
        else:
            temp = temp*np.array([matrix[0][i],matrix[1][i],matrix[2][i],matrix[3][i]]) + 0.2 * np.array([matrix[0][i],matrix[1][i],matrix[2][i],matrix[3][i]])
        #print(temp)
        thetaGuess = thetaGuess - temp
        # print(str(thetaGuess) + " + 'old'")
        # print(str(thetaGuess) + " + 'guess'")
        if (np.linalg.norm(thetaGuess - oldGuess) < tol) or (maxRuns == 1000):
            break
    else:
        maxRuns = maxRuns + 1
        if i == 129:
            i = 0
    # print(thetaGuess)
    return thetaGuess

thetaGuess4Reg = gradDescent4(matrix4, y) # Find parameters for theta using Gradient Descent.
print("True values for theta = " + str(thetaValues4))
print("Values for theta using GDReg = " + str(thetaGuess4))
for i in range(150):
    yValuesReg[i] = (thetaGuess4Reg[0] + thetaGuess4Reg[1] * X['x'][i] + thetaGuess4Reg[2] * X['x^2'][i] + thetaGuess4Reg[3] * X['x^3'][i]) + noiseValues[i]
print(" ")

#Computing the training error of Gradient of Descent of third order polynomial:
for i in range(90):
    total = 0
    E = (data['y'][i] - yValuesReg[i])**2
    total = total + E
trainingErrorReg = 0.5 * total
print("Training Error for Gradient Descent with regularization = " + str(trainingErrorReg))

#Computing the validation error of Gradient Descent of third order polynomial:
for i in range(30):
    total = 0
    E = (data['y'][90 + i] - yValuesReg[90 + i])**2
    total = total + E
validationErrorReg = 0.5 * total
print("Validation Error for Gradient Descent with regularization = " + str(validationErrorReg))

#Computing the test error of Gradient Descent:
for i in range(30):
    total = 0
    E = (data['y'][110 + i] - yValuesReg[110 + i])**2
    total = total + E
testErrorReg = 0.5 * total
print("Test Error for Gradient Descent with regularization = " + str(testErrorReg))

True values for theta = [0.75206039 0.33073176 0.48262031 0.54563629]
Values for theta using GDReg = [0.72234895 0.65362995 0.79008449 0.99912847]

Training Error for Gradient Descent with regularization = 0.080237268072602955
Validation Error for Gradient Descent with regularization = 0.020780190973927883
Test Error for Gradient Descent with regularization = 463.157968675616
```

iv. Compare your results of the three gradient descent based models, which model achieves the best final training error? Which model achieves the best validation error? Can you see any visible difference in the function approximation (fit of the data) by